

Statistical learning theory

Weekend 1, Friday late afternoon

Carlos Cotrini

Friday 4 September 2026

1

Section 1

Empirical risk minimisation

Vladimir Vapnik



Vladimir Vapnik



Second edition, 2006

Estimation of Dependences Based on Empirical Data. Russian original 1979, Kotz translation Springer 1982, second edition 2006.

Vladimir Vapnik



Vladimir Vapnik



Second edition, 2006

- He showed that catching spam, fitting a price and estimating a probability are **one problem**, scored three ways.

Estimation of Dependences Based on Empirical Data. Russian original 1979, Kotz translation Springer 1982, second edition 2006.

Vladimir Vapnik



Vladimir Vapnik



Second edition, 2006

- He showed that catching spam, fitting a price and estimating a probability are **one problem**, scored three ways.
- **Catching spam.** The score counts how often the answer is wrong.

Estimation of Dependences Based on Empirical Data. Russian original 1979, Kotz translation Springer 1982, second edition 2006.

Vladimir Vapnik



Vladimir Vapnik



Second edition, 2006

- He showed that catching spam, fitting a price and estimating a probability are **one problem**, scored three ways.
- **Catching spam.** The score counts how often the answer is wrong.
- **Fitting a price.** The score is the squared miss, the one this room has used all day.

Estimation of Dependences Based on Empirical Data. Russian original 1979, Kotz translation Springer 1982, second edition 2006.

Vladimir Vapnik



Vladimir Vapnik



Second edition, 2006

- He showed that catching spam, fitting a price and estimating a probability are **one problem**, scored three ways.
- **Catching spam.** The score counts how often the answer is wrong.
- **Fitting a price.** The score is the squared miss, the one this room has used all day.
- **Estimating a probability.** The score is how surprised the model is by what actually happened.

Estimation of Dependences Based on Empirical Data. Russian original 1979, Kotz translation Springer 1982, second edition 2006.

Vladimir Vapnik



Vladimir Vapnik



Second edition, 2006

- He showed that catching spam, fitting a price and estimating a probability are **one problem**, scored three ways.
- **Catching spam.** The score counts how often the answer is wrong.
- **Fitting a price.** The score is the squared miss, the one this room has used all day.
- **Estimating a probability.** The score is how surprised the model is by what actually happened.
- Three problems that had lived in different books. The score says which one you are doing, and what happens to the score afterwards never changes.


Estimation of Dependences Based on Empirical Data. Russian original 1979, Kotz translation Springer 1982, second edition 2006.

phenomenon



$x \longrightarrow y$


$y = f(x)$



④ Training algo $\Rightarrow$ estimator $\hat{f}$
 f minimizer $L(w)$

③ Loss function

② Model



lin model
neural network
⋮

① observations

(x_1, y_1) (x_2, y_2) ... (x_n, y_n)


unknown

$\hat{y}_1$ $\hat{y}_2$... $\hat{y}_n$

The number we want and the number we have

$\mathcal{M}$ is a **candidate model**: any rule that turns a row of data into an estimate.

Every case the world will hand us




$R(\mathcal{M})$, the average miss over all of them

Nobody can compute this number

The number we want and the number we have

$\mathcal{M}$ is a **candidate model**: any rule that turns a row of data into an estimate.

Every case the world will hand us



$R(\mathcal{M})$, the average miss over all of them

Nobody can compute this number

The n rows we collected




$\bar{I}(\mathcal{M})$, the average miss over those

The number we want and the number we have

$\mathcal{M}$ is a **candidate model**: any rule that turns a row of data into an estimate.

Every case the world will hand us



$R(\mathcal{M})$, the average miss over all of them

Nobody can compute this number

The n rows we collected




$\bar{I}(\mathcal{M})$, the average miss over those

Vapnik's name for it: **the empirical risk**

The number we want and the number we have

$\mathcal{M}$ is a **candidate model**: any rule that turns a row of data into an estimate.

Every case the world will hand us



$R(\mathcal{M})$, the average miss over all of them

Nobody can compute this number

The n rows we collected



$\bar{I}(\mathcal{M})$, the average miss over those

Vapnik's name for it: the empirical risk

Empirical risk minimisation: choose the $\mathcal{M}$ that makes $\bar{I}$ small, and take that as the answer to a question about R .


2

Section 2

The loop we have been running


Linear regression with gradient descent

Two o'clock, in one picture. Each box is a piece of data, and its shape is the shape of the data.




Linear regression with gradient descent

Two o'clock, in one picture. Each box is a piece of data, and its shape is the shape of the data.




Linear regression with gradient descent

Two o'clock, in one picture. Each box is a piece of data, and its shape is the shape of the data.




Linear regression with gradient descent


Two o'clock, in one picture. Each box is a piece of data, and its shape is the shape of the data.




A neural network




A neural network



A neural network




A neural network




Everything between the data and the estimate

The three weight blocks and the two tables between them are the model. Give the whole of it one name.




Everything between the data and the estimate

The three weight blocks and the two tables between them are the model. Give the whole of it one name.



Everything between the data and the estimate

The three weight blocks and the two tables between them are the model. Give the whole of it one name.



3


Section 3

The same loop, elsewhere

The loop, with the model as one box

The same picture, with $\mathcal{M}$ where the network was.


The $\square$ ve $\square$ ats are the ones from this morning, and the network is the one from this afternoon.



The loop, with the model as one box

The same picture, with $\mathcal{M}$ where the network was.


The $\square$ ve $\square$ ats are the ones from this morning, and the network is the one from this afternoon.



The loop, with the model as one box

The same picture, with $\mathcal{M}$ where the network was.


The $\square$ ve $\square$ ats are the ones from this morning, and the network is the one from this afternoon.



The loop, with the model as one box

The same picture, with $\mathcal{M}$ where the network was.


The $\square$ ve $\square$ ats are the ones from this morning, and the network is the one from this afternoon.



The loop, with the model as one box

The same picture, with $\mathcal{M}$ where the network was.

The $\nabla \bar{I}$ is the one from this morning, and the network is the one from this afternoon.




Large Language Model

Transformer 2017

$$\nabla l = (\dots)$$

②



① Data

y
upon
a
time
there
⋮

Once
upon
a
time
⋮

x_{11}	x_{12}	...	x_{1d}
x_{21}	x_{22}	...	x_{2d}
...			
x_{n1}			x_{nd}

X

upon
a
stepmother
fairy


l

l

Large language models

A row is text with its next token cut off. Every position of every document gives one.


The box holds a transformer. Vaswani and others, *Attention is all you need*, 2017.



Large language models

A row is text with its next token cut off. Every position of every document gives one.


The box holds a transformer. Vaswani and others, *Attention is all you need*, 2017.



Large language models

A row is text with its next token cut off. Every position of every document gives one.


The box holds a transformer. Vaswani and others, *Attention is all you need*, 2017.



Large language models

A row is text with its next token cut off. Every position of every document gives one.


The box holds a transformer. Vaswani and others, *Attention is all you need*, 2017.



Large language models

A row is text with its next token cut off. Every position of every document gives one.


The box holds a transformer. Vaswani and others, *Attention is all you need*, 2017.



Reasoning models

Start from a trained language model and keep going, on problems a machine can mark.


The verdict is a marker rather than a subtraction. DeepSeek-R1, *Nature* 645, 633 (2025).



Reasoning models

Start from a trained language model and keep going, on problems a machine can mark.


The verdict is a marker rather than a subtraction. DeepSeek-R1, *Nature* 645, 633 (2025).



Reasoning models

Start from a trained language model and keep going, on problems a machine can mark.


The verdict is a marker rather than a subtraction. DeepSeek-R1, *Nature* 645, 633 (2025).



Reasoning models

Start from a trained language model and keep going, on problems a machine can mark.


The verdict is a marker rather than a subtraction. DeepSeek-R1, *Nature* 645, 633 (2025).



Reasoning models

Start from a trained language model and keep going, on problems a machine can mark.


The verdict is a marker rather than a subtraction. DeepSeek-R1, *Nature* 645, 633 (2025).



Instruction tuning

GPT-3 was $\square$ nished in 2020 and it completed text. A smaller second pass made it answer.


12 725 training prompts, and its replies were preferred to a model 100 times larger. Ouyang and others, 2022.



Instruction tuning

GPT-3 was $\square$ nished in 2020 and it completed text. A smaller second pass made it answer.


12 725 training prompts, and its replies were preferred to a model 100 times larger. Ouyang and others, 2022.



Instruction tuning

GPT-3 was $\square$ nished in 2020 and it completed text. A smaller second pass made it answer.


12 725 training prompts, and its replies were preferred to a model 100 times larger. Ouyang and others, 2022.



Instruction tuning

GPT-3 was $\square$ nished in 2020 and it completed text. A smaller second pass made it answer.

12 725 training prompts, and its replies were preferred to a model 100 times larger. Ouyang and others, 2022.



Instruction tuning

GPT-3 was $\square$ nished in 2020 and it completed text. A smaller second pass made it answer.

12 725 training prompts, and its replies were preferred to a model 100 times larger. Ouyang and others, 2022.




Image generation

Add a measured amount of noise to a picture, then ask which noise was added.

The loss here really is the squared miss. Ho, Jain and Abbeel, 2020; Rombach and others, 2022.




Image generation

Add a measured amount of noise to a picture, then ask which noise was added.

The loss here really is the squared miss. Ho, Jain and Abbeel, 2020; Rombach and others, 2022.




Image generation

Add a measured amount of noise to a picture, then ask which noise was added.

The loss here really is the squared miss. Ho, Jain and Abbeel, 2020; Rombach and others, 2022.




Image generation

Add a measured amount of noise to a picture, then ask which noise was added.

The loss here really is the squared miss. Ho, Jain and Abbeel, 2020; Rombach and others, 2022.





Image generation

Add a measured amount of noise to a picture, then ask which noise was added.


The loss here really is the squared miss. Ho, Jain and Abbeel, 2020; Rombach and others, 2022.



Transcription

Every recording that already carries a transcript is a row. There are 680 000 hours.


Whisper. Radford and others, ICML 2023.



Transcription

Every recording that already carries a transcript is a row. There are 680 000 hours.


Whisper. Radford and others, ICML 2023.



Transcription

Every recording that already carries a transcript is a row. There are 680 000 hours.


Whisper. Radford and others, ICML 2023.



Transcription

Every recording that already carries a transcript is a row. There are 680 000 hours.


Whisper. Radford and others, ICML 2023.



Transcription

Every recording that already carries a transcript is a row. There are 680 000 hours.


Whisper. Radford and others, ICML 2023.



Ranking and recommendation

Every impression is a row: what was shown, to whom, and whether they took it.


The system that picks the next video. Covington, Adams and Sargin, RecSys 2016.



Ranking and recommendation

Every impression is a row: what was shown, to whom, and whether they took it.


The system that picks the next video. Covington, Adams and Sargin, RecSys 2016.



Ranking and recommendation

Every impression is a row: what was shown, to whom, and whether they took it.


The system that picks the next video. Covington, Adams and Sargin, RecSys 2016.



Ranking and recommendation

Every impression is a row: what was shown, to whom, and whether they took it.


The system that picks the next video. Covington, Adams and Sargin, RecSys 2016.



Ranking and recommendation

Every impression is a row: what was shown, to whom, and whether they took it.


The system that picks the next video. Covington, Adams and Sargin, RecSys 2016.



Fraud and credit decisions

Every settled transaction is a row, and its label arrives weeks later with the chargeback.


Dal Pozzolo and others, 2018, with Worldline. The tree ensemble is fitted to the gradient too.



Fraud and credit decisions

Every settled transaction is a row, and its label arrives weeks later with the chargeback.


Dal Pozzolo and others, 2018, with Worldline. The tree ensemble is fitted to the gradient too.



Fraud and credit decisions

Every settled transaction is a row, and its label arrives weeks later with the chargeback.


Dal Pozzolo and others, 2018, with Worldline. The tree ensemble is fitted to the gradient too.



Fraud and credit decisions

Every settled transaction is a row, and its label arrives weeks later with the chargeback.


Dal Pozzolo and others, 2018, with Worldline. The tree ensemble is fitted to the gradient too.



Fraud and credit decisions

Every settled transaction is a row, and its label arrives weeks later with the chargeback.


Dal Pozzolo and others, 2018, with Worldline. The tree ensemble is ∇ tted to the gradient too.



Screening for diabetic eye disease

In 2018 a regulator cleared a machine to read these photographs with no doctor behind it.


128 175 photographs, graded 3 to 7 times by 54 ophthalmologists. Gulshan and others, *JAMA*, 2016.



Screening for diabetic eye disease

In 2018 a regulator cleared a machine to read these photographs with no doctor behind it.


128 175 photographs, graded 3 to 7 times by 54 ophthalmologists. Gulshan and others, *JAMA*, 2016.



Screening for diabetic eye disease

In 2018 a regulator cleared a machine to read these photographs with no doctor behind it.


128 175 photographs, graded 3 to 7 times by 54 ophthalmologists. Gulshan and others, *JAMA*, 2016.



Screening for diabetic eye disease

In 2018 a regulator cleared a machine to read these photographs with no doctor behind it.


128 175 photographs, graded 3 to 7 times by 54 ophthalmologists. Gulshan and others, *JAMA*, 2016.



Screening for diabetic eye disease

In 2018 a regulator cleared a machine to read these photographs with no doctor behind it.

128 175 photographs, graded 3 to 7 times by 54 ophthalmologists. Gulshan and others, *JAMA*, 2016.



References: Vapnik and the language models

Vladimir N. Vapnik. *Estimation of Dependences Based on Empirical Data*. Translated by Samuel Kotz. Springer, 1982; second edition 2006.

doi.org/10.1007/0-387-34239-7

Ashish Vaswani and others. Attention is all you need. *Advances in Neural Information Processing Systems* 30, 2017.

arxiv.org/abs/1706.03762

Long Ouyang and others. Training language models to follow instructions with human feedback, 2022.

arxiv.org/abs/2203.02155

DeepSeek-AI. DeepSeek-R1 incentivizes reasoning in LLMs through reinforcement learning. *Nature* 645, 633 to 638, 2025.

arxiv.org/abs/2501.12948

References: the other applications

Jonathan Ho, Ajay Jain and Pieter Abbeel. Denoising diffusion probabilistic models. *NeurIPS* 33, 2020.

arxiv.org/abs/2006.11239

Robin Rombach and others. High-resolution image synthesis with latent diffusion models. *CVPR*, 2022.

arxiv.org/abs/2112.10752

Alec Radford and others. Robust speech recognition via large-scale weak supervision. *ICML*, PMLR 202, 28492, 2023.

arxiv.org/abs/2212.04356

Paul Covington, Jay Adams and Emre Sargin. Deep neural networks for YouTube recommendations. *RecSys*, 191 to 198, 2016.

doi.org/10.1145/2959100.2959190

Andrea Dal Pozzolo and others. Credit card fraud detection: a realistic modeling and a novel learning strategy. *IEEE Transactions on Neural Networks and Learning Systems* 29(8), 3784 to 3797, 2018.

doi.org/10.1109/TNNLS.2017.2736643

Tianqi Chen and Carlos Guestrin. XGBoost: a scalable tree boosting system. *KDD*, 2016. arxiv.org/abs/1603.02754

Varun Gulshan and others. Development and validation of a deep learning algorithm for detection of diabetic retinopathy in retinal fundus photographs. *JAMA* 316(22), 2402 to 2410, 2016. doi.org/10.1001/jama.2016.17216