

Training neural networks with PyTorch

Weekend 1, Saturday morning

Carlos Cotrini

ETH Zurich

Saturday 5 September 2026

1

Section 1

What Friday's loop costs

Which iris is it?

Three species, four measurements, one guess. That is the whole problem, and the rest of the hour is how a computer does it.



Setosa



Versicolor



Virginica

Which iris is it?

Three species, four measurements, one guess. That is the whole problem, and the rest of the hour is how a computer does it.



Setosa



Versicolor



Virginica

One flower, measured

6.1	2.8	4.7	1.2
-----	-----	-----	-----

Sepal and petal, each a length and a width

Which iris is it?

Three species, four measurements, one guess. That is the whole problem, and the rest of the hour is how a computer does it.



Setosa



Versicolor



Virginica

One flower, measured



Sepal and petal, each a length and a width



The network

Which iris is it?

Three species, four measurements, one guess. That is the whole problem, and the rest of the hour is how a computer does it.




Setosa



Versicolor



Virginica



Which iris is it?

Three species, four measurements, one guess. That is the whole problem, and the rest of the hour is how a computer does it.



Setosa



Versicolor



Virginica

One flower, measured



Sepal and petal, each a length and a width

It really is a versicolor, and this same flower comes back at the end of the hour.




The network

Versicolor

the network's answer


The picture from Friday afternoon

Friday afternoon, unchanged. Each box is a piece of data, and its shape is the shape of the data.




The picture from Friday afternoon

Friday afternoon, unchanged. Each box is a piece of data, and its shape is the shape of the data.




The picture from Friday afternoon

Friday afternoon, unchanged. Each box is a piece of data, and its shape is the shape of the data.




The picture from Friday afternoon

Friday afternoon, unchanged. Each box is a piece of data, and its shape is the shape of the data.




The picture from Friday afternoon

Friday afternoon, unchanged. Each box is a piece of data, and its shape is the shape of the data.




The numbers this network holds

Four measurements in, sixteen hidden values in the middle, three species scores out.




The numbers this network holds

Four measurements in, sixteen hidden values in the middle, three species scores out.




The numbers this network holds

Four measurements in, sixteen hidden values in the middle, three species scores out.




The numbers this network holds

Four measurements in, sixteen hidden values in the middle, three species scores out.



The numbers this network holds

Four measurements in, sixteen hidden values in the middle, three species scores out.



Four objects, and what each one does


Everything this hour does is these four things, in a circle, two hundred times.

A tensor

A block of numbers. The measurements are one, and so is every weight.

Four objects, and what each one does

Everything this hour does is these four things, in a circle, two hundred times.




A block of numbers. The measurements are one, and so is every weight.

The network. It holds the weights and turns one tensor into another.


Four objects, and what each one does

Everything this hour does is these four things, in a circle, two hundred times.




Four objects, and what each one does

Everything this hour does is these four things, in a circle, two hundred times.



Four objects, and what each one does

Everything this hour does is these four things, in a circle, two hundred times.



What replaces what

The data, as a block of numbers

`torch.tensor(...)`

What replaces what

The data, as a block of numbers

`torch.tensor(...)`

The network of boxes

`nn.Linear`, `nn.ReLU`, `nn.Module`

What replaces what

The data, as a block of numbers

`torch.tensor(...)`

The network of boxes

`nn.Linear`, `nn.ReLU`, `nn.Module`

The verdict, how wrong we are

`nn.CrossEntropyLoss()`

What replaces what

The data, as a block of numbers

`torch.tensor(...)`

The network of boxes

`nn.Linear, nn.ReLU, nn.Module`

The verdict, how wrong we are

`nn.CrossEntropyLoss()`

The tilt, measured in a window

`loss.backward()`

What replaces what

The data, as a block of numbers

`torch.tensor(...)`

The network of boxes

`nn.Linear, nn.ReLU, nn.Module`

The verdict, how wrong we are

`nn.CrossEntropyLoss()`

The tilt, measured in a window

`loss.backward()`

$w = w - \alpha * g$

`optimizer.step()`

What replaces what

The data, as a block of numbers

`torch.tensor(...)`

The network of boxes

`nn.Linear, nn.ReLU, nn.Module`

The verdict, how wrong we are

`nn.CrossEntropyLoss()`

The tilt, measured in a window

`loss.backward()`

`w = w - alpha * g`

`optimizer.step()`

The step size, α

`lr=0.05`

2

Section 2

The gradient and the update


The gradient, for every weight at once

Friday, 11:00. Measuring the tilt of one weight.

$$a = L(w - \text{eps}/2)$$

$$b = L(w + \text{eps}/2)$$

$$g = (b - a) / \text{eps}$$




The gradient, for every weight at once

Friday, 11:00. Measuring the tilt of one weight.

$$a = L(w - \text{eps}/2)$$

$$b = L(w + \text{eps}/2)$$

$$g = (b - a) / \text{eps}$$



The gradient, for every weight at once


Friday, 11:00. Measuring the tilt of one weight.

$$a = L(w - \text{eps}/2)$$

$$b = L(w + \text{eps}/2)$$

$$g = (b - a) / \text{eps}$$

`loss.backward()`



The gradient, for every weight at once

Friday, 11:00. Measuring the tilt of one weight.


~~$$a = L(w - \text{eps}/2)$$~~

~~$$b = L(w + \text{eps}/2)$$~~

~~$$g = (b - a) / \text{eps}$$~~

`loss.backward()`

One call, and every one of the 131 numbers has its own tilt waiting on it.



The update rule, in one call

The step Friday took by hand, once for each weight.


$$w = w - \alpha * g$$

One weight

The update rule, in one call

The step Friday took by hand, once for each weight.

$$w = w - \alpha * g$$




The update rule, in one call

The step Friday took by hand, once for each weight.

$$w = w - \text{alpha} * g$$

```
optimizer.step()
```




The update rule, in one call

The step Friday took by hand, once for each weight.

$$\underline{w = w - \text{alpha} * g}$$

`optimizer.step()`



The update rule, in one call

The step Friday took by hand, once for each weight.


~~$w = w - \alpha * g$~~

`optimizer.step()`

One call moves all 131.

131 numbers

One weight



3

Section 3

The notebook

The data


The Iris measurements, and the two blocks every later frame is drawn from.



```
import torch
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
```


The data

The Iris measurements, and the two blocks every later frame is drawn from.




```
import torch
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split

iris = load_iris()
X = iris.data
y = iris.target
```

The data

The Iris measurements, and the two blocks every later frame is drawn from.



X Sepal length
Sepal width
Petal length
Petal width

y Setosa
Versicolor
Virginica


```
import torch
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split

iris = load_iris()
X = iris.data
y = iris.target

print("Features shape:", X.shape)
print("Labels shape: ", y.shape)
print("Classes:      ", iris.target_names)
```


The data

The Iris measurements, and the two blocks every later frame is drawn from.



The data

The Iris measurements, and the two blocks every later frame is drawn from.




Thirty of the flowers are held back so the score at the end is a fair one.
Saturday 10:30 says why that matters.

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42  
)  
  
X_train = torch.tensor(X_train, dtype=torch.float32)  
X_test  = torch.tensor(X_test,  dtype=torch.float32)  
y_train = torch.tensor(y_train, dtype=torch.long)  
y_test  = torch.tensor(y_test,  dtype=torch.long)
```

The network, as blocks


The same picture as yesterday, with the notebook's own widths.

```
import torch.nn as nn
```



The network, as blocks

The same picture as yesterday, with the notebook's own widths.




```
import torch.nn as nn
```

```
self.layer1 = nn.Linear(4, 16)
```

The network, as blocks

The same picture as yesterday, with the notebook's own widths.




```
import torch.nn as nn
```

```
self.layer1 = nn.Linear(4, 16)
```

```
self.relu = nn.ReLU()
```

The network, as blocks

The same picture as yesterday, with the notebook's own widths.



```
import torch.nn as nn
```


```
self.layer1 = nn.Linear(4, 16)
```

```
self.relu = nn.ReLU()
```

```
self.layer2 = nn.Linear(16, 3)
```


The network, as blocks

The same picture as yesterday, with the notebook's own widths.



```
import torch.nn as nn
```

```
self.layer1 = nn.Linear(4, 16)
self.relu   = nn.ReLU()
self.layer2 = nn.Linear(16, 3)
```

The network, as a class

```
class IrisNet(nn.Module):  
    def __init__(self):  
        super().__init__()  
  
    def forward(self, x):
```

The parts {

The wiring {

The network, as a class

The parts {

```
class IrisNet(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.layer1 = nn.Linear(4, 16)  
        self.relu    = nn.ReLU()  
        self.layer2 = nn.Linear(16, 3)
```

The wiring {

```
    def forward(self, x):
```

The network, as a class

The parts {

```
class IrisNet(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.layer1 = nn.Linear(4, 16)  
        self.relu    = nn.ReLU()  
        self.layer2 = nn.Linear(16, 3)
```

1
2
3 { The wiring

```
    def forward(self, x):  
        x = self.layer1(x)  
        x = self.relu(x)  
        x = self.layer2(x)  
        return x
```

The network, as a class

```
class IrisNet(nn.Module):
```

```
model = IrisNet()
```

The parts

```
def __init__(self):
```

```
    super().__init__()
```

```
    self.layer1 = nn.Linear(4, 16)
```

```
    self.relu = nn.ReLU()
```

```
    self.layer2 = nn.Linear(16, 3)
```

1

2

3

The wiring

```
def forward(self, x):
```

```
    x = self.layer1(x)
```

```
    x = self.relu(x)
```

```
    x = self.layer2(x)
```

```
    return x
```

The network, as a class

```
class IrisNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.layer1 = nn.Linear(4, 16)
        self.relu = nn.ReLU()
        self.layer2 = nn.Linear(16, 3)

    def forward(self, x):
        x = self.layer1(x)
        x = self.relu(x)
        x = self.layer2(x)
        return x

model = IrisNet()
print(model)
```

The parts {

1
2
3 The wiring {

```
IrisNet(
  (layer1): Linear(in_features=4, out_features=16, bias=True)
  (relu): ReLU()
  (layer2): Linear(in_features=16, out_features=3, bias=True)
)
```


The network, drawn as neurons

One circle for every column of every weight block.

Sepal length ○
Sepal width ○
Petal length ○
Petal width ○


The network, drawn as neurons

One circle for every column of every weight block.




The network, drawn as neurons

One circle for every column of every weight block.



The network, drawn as neurons


One circle for every column of every weight block.




The loss and the optimizer

Friday's miss measured a number. This one scores a choice among three species.


The true species



Setosa




Versicolor



Virginica

The loss and the optimizer


Friday's miss measured a number. This one scores a choice among three species.



```
loss_fn = nn.CrossEntropyLoss()
```

The loss and the optimizer


Friday's miss measured a number. This one scores a choice among three species.



```
loss_fn    = nn.CrossEntropyLoss()  
optimizer = torch.optim.SGD(model.parameters(), lr=0.05)
```

The loss and the optimizer

Friday's miss measured a number. This one scores a choice among three species.




The network's 131 numbers.

```
loss_fn = nn.CrossEntropyLoss()  
optimizer = torch.optim.SGD(model.parameters(), lr=0.05)
```

The loss and the optimizer

Friday's miss measured a number. This one scores a choice among three species.



The training loop


Data

```
num_epochs = 200
```

```
loss_history = []
```

```
for epoch in range(num_epochs):
```


The training loop




```
num_epochs = 200
```

```
loss_history = []
```

```
for epoch in range(num_epochs):  
    predictions = model(X_train)
```

The training loop




```
num_epochs = 200
```

```
loss_history = []
```

```
for epoch in range(num_epochs):  
    predictions = model(X_train)  
    loss = loss_fn(predictions, y_train)
```


The training loop



```
num_epochs = 200  
loss_history = []
```

```
for epoch in range(num_epochs):  
    predictions = model(X_train)  
    loss = loss_fn(predictions, y_train)  
    optimizer.zero_grad()  
    loss.backward()
```


The training loop



```
num_epochs = 200  
loss_history = []
```

```
for epoch in range(num_epochs):  
    predictions = model(X_train)  
    loss = loss_fn(predictions, y_train)  
    optimizer.zero_grad()  
    loss.backward()  
    optimizer.step()
```


The training loop



```
num_epochs = 200  
loss_history = []
```

```
for epoch in range(num_epochs):  
    predictions = model(X_train)  
    loss = loss_fn(predictions, y_train)  
    optimizer.zero_grad()  
    loss.backward()  
    optimizer.step()  
    loss_history.append(loss.item())
```

The training loop




```
num_epochs = 200  
loss_history = []
```

```
for epoch in range(num_epochs):  
    predictions = model(X_train)  
    loss = loss_fn(predictions, y_train)  
    optimizer.zero_grad()  
    loss.backward()  
    optimizer.step()  
    loss_history.append(loss.item())
```


The learning curve

One seeded run. Your own run will differ.



The learning curve


One seeded run. Your own run will differ.



```
plt.plot(loss_history)
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.show()
```


The learning curve

One seeded run. Your own run will differ.




```
plt.plot(loss_history)
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.show()
```

Epoch	50 / 200		Loss:	0.4802
Epoch	100 / 200		Loss:	0.3505
Epoch	150 / 200		Loss:	0.2665
Epoch	200 / 200		Loss:	0.2082

The learning curve

One seeded run. Your own run will differ.



```
plt.plot(loss_history)
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.show()
```

Epoch	50 / 200		Loss:	0.4802
Epoch	100 / 200		Loss:	0.3505
Epoch	150 / 200		Loss:	0.2665
Epoch	200 / 200		Loss:	0.2082

The test set

One seeded run. Your own run will differ.

```
with torch.no_grad():  
    test_scores = model(X_test)  
    predicted   = test_scores.argmax(1)
```



The test set

One seeded run. Your own run will differ.

```
with torch.no_grad():  
    test_scores = model(X_test)  
    predicted    = test_scores.argmax(1)
```

The call `argmax(1)` keeps the species with the largest score.



The test set

One seeded run. Your own run will differ.

```
with torch.no_grad():  
    test_scores = model(X_test)  
    predicted    = test_scores.argmax(1)
```

The call `argmax(1)` keeps the species with the largest score.

One of the thirty: a virginica
read as a versicolor.



The test set

One seeded run. Your own run will differ.

```
with torch.no_grad():  
    test_scores = model(X_test)  
    predicted    = test_scores.argmax(1)
```

The call `argmax(1)` keeps the species with the largest score.

```
correct  = (predicted == y_test).sum().item()  
total    = y_test.shape[0]  
accuracy = correct / total * 100
```

One of the thirty: a virginica
read as a versicolor.



The test set

One seeded run. Your own run will differ.

```
with torch.no_grad():  
    test_scores = model(X_test)  
    predicted    = test_scores.argmax(1)
```

The call `argmax(1)` keeps the species with the largest score.

```
correct  = (predicted == y_test).sum().item()  
total    = y_test.shape[0]  
accuracy = correct / total * 100
```

```
print(f"Correct: {correct} / {total}")  
print(f"Accuracy: {accuracy:.1f}%")
```

One of the thirty: a virginica
read as a versicolor.



Correct: 29 / 30

Accuracy: 96.7%

One prediction up close

One seeded run. Your own run will differ.

```
flower = X_test[0].unsqueeze(0)
true_label = y_test[0].item()
```

One test flower			
6.1	2.8	4.7	1.2
Sepal length	Sepal width	Petal length	Petal width

One prediction up close

One seeded run. Your own run will differ.

```
flower = X_test[0].unsqueeze(0)
true_label = y_test[0].item()
```

```
with torch.no_grad():
    scores = model(flower)
```

One test flower			
6.1	2.8	4.7	1.2
Sepal length	Sepal width	Petal length	Petal width
Raw scores out of layer2			
-2.6937	+1.9177	+1.3280	
Setosa	Versicolor	Virginica	



One prediction up close

One seeded run. Your own run will differ.

```
flower = X_test[0].unsqueeze(0)
true_label = y_test[0].item()
```

```
with torch.no_grad():
    scores = model(flower)
```

```
probs = torch.softmax(scores, dim=1)[0]
```





One prediction up close

One seeded run. Your own run will differ.

```
flower = X_test[0].unsqueeze(0)
true_label = y_test[0].item()
```

```
with torch.no_grad():
    scores = model(flower)
```

```
probs = torch.softmax(scores, dim=1)[0]
```




One prediction up close

One seeded run. Your own run will differ.


```
flower = X_test[0].unsqueeze(0)
true_label = y_test[0].item()
```

```
with torch.no_grad():
    scores = model(flower)
```

```
probs = torch.softmax(scores, dim=1)[0]
```



The network returns these three raw scores. The softmax runs here, once, after training, so the numbers can be read as probabilities.




4

Section 4


The knobs

Four things to change




- 1 Change the step: try 0.1 or 0.01.

Four things to change




- 1 Change the step: try 0.1 or 0.01.
- 2 Change the hidden layer: 16 becomes 4 or 64.

Four things to change




- ① Change the step: try 0.1 or 0.01.
- ② Change the hidden layer: 16 becomes 4 or 64.
- ③ Add a second hidden layer.

Four things to change



- ① Change the step: try 0.1 or 0.01.
- ② Change the hidden layer: 16 becomes 4 or 64.
- ③ Add a second hidden layer.
- ④ Train for more epochs, or fewer.

Four things to change



- 1 Change the step: try 0.1 or 0.01.
- 2 Change the hidden layer: 16 becomes 4 or 64.
- 3 Add a second hidden layer.
- 4 Train for more epochs, or fewer.

What replaced what

The data, as a block of numbers

The network of boxes

The verdict, how wrong we are

The tilt, measured in a window

$$w = w - \alpha * g$$

The step size, α

What replaced what

The data, as a block of numbers

- `torch.tensor(...)`

The network of boxes

- `nn.Linear`, `nn.ReLU`, `nn.Module`

The verdict, how wrong we are

- `nn.CrossEntropyLoss()`

The tilt, measured in a window

- `loss.backward()`

$w = w - \alpha * g$

- `optimizer.step()`

The step size, α

- `lr=0.05`