

Agentic AI, part 2

Agents

Carlos Cotrini

Weekend 2


1

Section 1

Agents and the loop


The agent and its environment

Sutton and Barto, Reinforcement Learning, 2018




The agent and its environment

Sutton and Barto, Reinforcement Learning, 2018




The agent and its environment

Sutton and Barto, Reinforcement Learning, 2018




The agent and its environment

Sutton and Barto, Reinforcement Learning, 2018




The agent and its environment

Sutton and Barto, Reinforcement Learning, 2018



The agent and its environment

Sutton and Barto, Reinforcement Learning, 2018



Environment


Observations

Actions

Reward

The agent and its environment

Sutton and Barto, Reinforcement Learning, 2018



The reward is always computed outside the agent


The spy game

The tax agent

Environment	The 3 by 3 base
Observations	What the game prints after each move
Actions	Three tools, and one scan the harness runs
Reward	None is computed

The agent and its environment


Sutton and Barto, Reinforcement Learning, 2018



	The spy game	The tax agent
Environment	The 3 by 3 base	Documents, a form, a mail account
Observations	What the game prints after each move	What a tool returns
Actions	Three tools, and one scan the harness runs	Read, extract, fill, send
Reward	None is computed	None is computed

The harness

The program around the model. The model only writes text; the harness turns that text into actions and feeds back what happened.




*Anthropic, Claude API documentation,
tool use, 2026*

*Young, Effective harnesses for
long-running agents, Anthropic, 2025*

The harness

The program around the model. The model only writes text; the harness turns that text into actions and feeds back what happened.




- 1 Builds the prompt: the tools and the transcript

*Anthropic, Claude API documentation,
tool use, 2026*

*Young, Effective harnesses for
long-running agents, Anthropic, 2025*

The harness

The program around the model. The model only writes text; the harness turns that text into actions and feeds back what happened.




- ① Builds the prompt: the tools and the transcript
- ② Reads the reply and finds the tool call

*Anthropic, Claude API documentation,
tool use, 2026*

*Young, Effective harnesses for
long-running agents, Anthropic, 2025*

The harness

The program around the model. The model only writes text; the harness turns that text into actions and feeds back what happened.




- ① Builds the prompt: the tools and the transcript
- ② Reads the reply and finds the tool call
- ③ Runs that tool on the world

*Anthropic, Claude API documentation,
tool use, 2026*

*Young, Effective harnesses for
long-running agents, Anthropic, 2025*

The harness

The program around the model. The model only writes text; the harness turns that text into actions and feeds back what happened.




- ① Builds the prompt: the tools and the transcript
- ② Reads the reply and finds the tool call
- ③ Runs that tool on the world
- ④ Appends the result, then calls the model again or stops

*Anthropic, Claude API documentation,
tool use, 2026*

*Young, Effective harnesses for
long-running agents, Anthropic, 2025*

The harness

The program around the model. The model only writes text; the harness turns that text into actions and feeds back what happened.



In the spy game: `run_agent()`

- ① Builds the prompt: the tools and the transcript
- ② Reads the reply and finds the tool call
- ③ Runs that tool on the world
- ④ Appends the result, then calls the model again or stops

*Anthropic, Claude API documentation,
tool use, 2026*

*Young, Effective harnesses for
long-running agents, Anthropic, 2025*

2

Section 2

ReAct, played by hand

The mission, and two ways to play it

Three dossiers or more, thirty turns, three tools.



Actions only

1

2

3

4

5

6

A thought, then an action

	0	1	2
0	Jungle	Open	Cache
1	Open	Dropout	Open
2	● Start	Vapnik	Robot

The mission, and two ways to play it

Three dossiers or more, thirty turns, three tools.



Actions only

1

2

3

4

5

6

A thought, then an action

	0	1	2
0	Jungle	Open	Cache
1	Open	Dropout	Open
2	● Start	Vapnik	Robot

Jungle (0,0):
a Flamethrower
Cache (0,2): 1 dossier
Robot (2,2):
the Cryo-Sentinel

The mission, and two ways to play it

Three dossiers or more, thirty turns, three tools.



Actions only

1

2

3

4

5

6

move talk collect

scan, run by the
harness every turn

A thought, then an action

	0	1	2
0	Jungle	Open	Cache
1	Open	Dropout	Open
2	Start	Vapnik	Robot

Jungle (0,0):
a Flamethrower
Cache (0,2): 1 dossier
Robot (2,2):
the Cryo-Sentinel

The mission, and two ways to play it

Three dossiers or more, thirty turns, three tools.



Actions only

- 1 Move east to (2,1), and meet Dr. Vapnik
- 2 Talk, and receive the USB Drive
- 3 Move north to (1,1), and meet Agent Dropout
- 4 Talk, and the harness drops the message
- 5 Talk, and the harness drops the message
- 6 Talk, and the harness drops the message

move talk collect

scan, run by the
harness every turn

A thought, then an action

	0	1	2
0	Jungle	Open	Cache
1	Open	Dropout	Open
2	Start	Vapnik	Robot

Jungle (0,0):
a Flamethrower
Cache (0,2): 1 dossier
Robot (2,2):
the Cryo-Sentinel

The mission, and two ways to play it

Three dossiers or more, thirty turns, three tools.



Actions only

- 1 Move east to (2,1), and meet Dr. Vapnik
- 2 Talk, and receive the USB Drive
- 3 Move north to (1,1), and meet Agent Dropout
- 4 Talk, and the harness drops the message
- 5 Talk, and the harness drops the message
- 6 Talk, and the harness drops the message

And 21 more. Won on the 27th action.

move talk collect

scan, run by the
harness every turn

A thought, then an action

	0	1	2
0	Jungle	Open	Cache
1	Open	Dropout	Open
2	Start	Vapnik	Robot

Jungle (0,0):
a Flamethrower
Cache (0,2): 1 dossier
Robot (2,2):
the Cryo-Sentinel

The mission, and two ways to play it

Three dossiers or more, thirty turns, three tools.



Actions only


- 1 Move east to (2,1), and meet Dr. Vapnik
- 2 Talk, and receive the USB Drive
- 3 Move north to (1,1), and meet Agent Dropout
- 4 Talk, and the harness drops the message
- 5 Talk, and the harness drops the message
- 6 Talk, and the harness drops the message

And 21 more. Won on the 27th action.

move talk collect

scan, run by the
harness every turn

A thought, then an action



Jungle (0,0):
a Flamethrower
Cache (0,2): 1 dossier
Robot (2,2):
the Cryo-Sentinel

The mission, and two ways to play it

Three dossiers or more, thirty turns, three tools.

Actions only


- 1 Move east to (2,1), and meet Dr. Vapnik
- 2 Talk, and receive the USB Drive
- 3 Move north to (1,1), and meet Agent Dropout
- 4 Talk, and the harness drops the message
- 5 Talk, and the harness drops the message
- 6 Talk, and the harness drops the message

And 21 more. Won on the 27th action.

{ move talk collect }

scan, run by the
harness every turn

A thought, then an action



Jungle (0,0):
a Flamethrower
Cache (0,2): 1 dossier
Robot (2,2):
the Cryo-Sentinel

Running total: 1,
then 2, and the turns
ran out

The mission, and two ways to play it

Three dossiers or more, thirty turns, three tools.

Actions only

Same map, same tools, same model

- 1 Move east to (2,1), and meet Dr. Vapnik
- 2 Talk, and receive the USB Drive
- 3 Move north to (1,1), and meet Agent Dropout
- 4 Talk, and the harness drops the message
- 5 Talk, and the harness drops the message
- 6 Talk, and the harness drops the message


And 21 more. Won on the 27th action.

{ move talk collect }

scan, run by the
harness every turn

A thought, then an action

One line added to the prompt



Jungle (0,0):
a Flamethrower
Cache (0,2): 1 dossier
Robot (2,2):
the Cryo-Sentinel

Running total: 1,
then 2, and the turns
ran out

The mission, and two ways to play it

Three dossiers or more, thirty turns, three tools.

Actions only

Same map, same tools, same model

- 1 Move east to (2,1), and meet Dr. Vapnik
- 2 Talk, and receive the USB Drive
- 3 Move north to (1,1), and meet Agent Dropout
- 4 Talk, and the harness drops the message
- 5 Talk, and the harness drops the message
- 6 Talk, and the harness drops the message


And 21 more. Won on the 27th action.

move talk collect

scan, run by the harness every turn

A thought, then an action

One line added to the prompt




Jungle (0,0):
a Flamethrower
Cache (0,2): 1 dossier
Robot (2,2):
the Cryo-Sentinel

Running total: 1,
then 2, and the turns
ran out

gemini-2.5-flash,
16 September 2026.
Without a thought
3 of 5 runs won.
With a thought 2 of 5.


Thought, action, observation



-  A call to the model
-  Something that is not a call to the model
-  The tools the model may choose from

Yao and others, ReAct, ICLR 2023

Thought, action, observation



A call to the model



Something that is not a call to the model




The tools the model may choose from



The stopping test


Yao and others, ReAct, ICLR 2023

Thought, action, observation




Yao and others, ReAct, ICLR 2023

Thought, action, observation




Yao and others, ReAct, ICLR 2023

Thought, action, observation



Thought, action, observation




- The model writes a thought. Nothing outside changes.
- It writes one action. The harness runs it.
- What comes back is appended, and the model is called again.

```
while not mission_over(operative):  
    observation = look_around(operative, world)  
    action = decide(briefing, transcript, observation)  
    result = execute(action, tools)  
    remember(transcript, observation, action, result)
```

- A call to the model
- Something that is not a call to the model
- The tools the model may choose from
- The stopping test

Yao and others, ReAct, ICLR 2023

Thought, action, observation



- The model writes a thought. Nothing outside changes.
- It writes one action. The harness runs it.
- What comes back is appended, and the model is called again.

```
while not mission_over(operative):  
    observation = look_around(operative, world)  
    action = decide(briefing, transcript, observation)  
    result = execute(action, tools)  
    remember(transcript, observation, action, result)
```

- A call to the model
- Something that is not a call to the model
- The tools the model may choose from
- The stopping test

The stopping test is in the loop's own condition. It is not a step counter. *Yao and others, ReAct, ICLR 2023*

The transcript the model is called with



Written once, by a person

You are a spy agent on a 3x3 grid.
Collect 3 dossiers in 30 turns.
Tools (use exactly one per turn):
 TOOL: move(direction="north")
 TOOL: talk(message="hello")
 TOOL: collect()
Respond with exactly one TOOL: line.
Write one THOUGHT: line, then
one TOOL: line.

Yao and others, ReAct, ICLR 2023
Rajasekaran and others, Effective
context engineering, Anthropic, 2025

The transcript the model is called with

The transcript



Written once, by a person

```
You are a spy agent on a 3x3 grid.  
Collect 3 dossiers in 30 turns.  
Tools (use exactly one per turn):  
  TOOL: move(direction="north")  
  TOOL: talk(message="hello")  
  TOOL: collect()  
Respond with exactly one TOOL: line.  
Write one THOUGHT: line, then  
one TOOL: line.
```

Yao and others, ReAct, ICLR 2023
Rajasekaran and others, Effective
context engineering, Anthropic, 2025

The transcript the model is called with

Observation

The transcript

[the mission briefing, 23 lines]

You are at position (2, 0).

East: Informant *[the scan's other 4 lines]*

The harness pasted the briefing here



Written once, by a person

You are a spy agent on a 3x3 grid.
Collect 3 dossiers in 30 turns.

Tools (use exactly one per turn):

TOOL: move(direction="north")

TOOL: talk(message="hello")

TOOL: collect()

Respond with exactly one TOOL: line.

Write one THOUGHT: line, then
one TOOL: line.

*Yao and others, ReAct, ICLR 2023
Rajasekaran and others, Effective
context engineering, Anthropic, 2025*

The transcript the model is called with

Observation

The transcript

[the mission briefing, 23 lines]

You are at position (2, 0).

East: Informant *[the scan's other 4 lines]*

The harness pasted the briefing here

Written once, by a person

You are a spy agent on a 3x3 grid.

Collect 3 dossiers in 30 turns.

Tools (use exactly one per turn):

TOOL: move(direction="north")

TOOL: talk(message="hello")

TOOL: collect()

Respond with exactly one TOOL: line.

Write one THOUGHT: line, then

one TOOL: line.

What the model wrote this step

[no thought on this turn]

Yao and others, ReAct, ICLR 2023
Rajasekaran and others, Effective
context engineering, Anthropic, 2025

The transcript the model is called with

Observation

The transcript

[the mission briefing, 23 lines]

You are at position (2, 0).

East: Informant *[the scan's other 4 lines]*

The harness pasted the briefing here

Written once, by a person

You are a spy agent on a 3x3 grid.

Collect 3 dossiers in 30 turns.

Tools (use exactly one per turn):

TOOL: move(direction="north")

TOOL: talk(message="hello")

TOOL: collect()

Respond with exactly one TOOL: line.

Write one THOUGHT: line, then


one TOOL: line.

What the model wrote this step

TOOL: move(direction="east")

Yao and others, ReAct, ICLR 2023
Rajasekaran and others, Effective
context engineering, Anthropic, 2025

The transcript the model is called with




Written once, by a person

You are a spy agent on a 3x3 grid.
Collect 3 dossiers in 30 turns.
Tools (use exactly one per turn):
TOOL: move(direction="north")
TOOL: talk(message="hello")
TOOL: collect()
Respond with exactly one TOOL: line.
Write one THOUGHT: line, then
one TOOL: line.

What the model wrote this step

Yao and others, ReAct, ICLR 2023
Rajasekaran and others, Effective
context engineering, Anthropic, 2025

The transcript the model is called with



Written once, by a person


You are a spy agent on a 3x3 grid.
Collect 3 dossiers in 30 turns.
Tools (use exactly one per turn):
TOOL: move(direction="north")
TOOL: talk(message="hello")
TOOL: collect()
Respond with exactly one TOOL: line.
Write one THOUGHT: line, then
one TOOL: line.

What the model wrote this step

TOOL: talk(message="Dr. Vapnik, I
seek the global minimum. Do you have
any jobs for an agent like me?")

Yao and others, ReAct, ICLR 2023
Rajasekaran and others, Effective
context engineering, Anthropic, 2025

The transcript the model is called with



Written once, by a person


You are a spy agent on a 3x3 grid.
Collect 3 dossiers in 30 turns.
Tools (use exactly one per turn):
TOOL: move(direction="north")
TOOL: talk(message="hello")
TOOL: collect()
Respond with exactly one TOOL: line.
Write one THOUGHT: line, then
one TOOL: line.

What the model wrote this step

TOOL: talk(message="Dr. Vapnik, I
seek the global minimum. Do you have
any jobs for an agent like me?")

Yao and others, ReAct, ICLR 2023
Rajasekaran and others, Effective
context engineering, Anthropic, 2025

The transcript the model is called with



Written once, by a person


You are a spy agent on a 3x3 grid.
Collect 3 dossiers in 30 turns.
Tools (use exactly one per turn):
TOOL: move(direction="north")
TOOL: talk(message="hello")
TOOL: collect()
Respond with exactly one TOOL: line.
Write one THOUGHT: line, then
one TOOL: line.

What the model wrote this step

TOOL: talk(message="Dr. Vapnik, I
seek the global minimum. Do you have
any jobs for an agent like me?")

Yao and others, ReAct, ICLR 2023
Rajasekaran and others, Effective
context engineering, Anthropic, 2025

The transcript the model is called with



Written once, by a person

You are a spy agent on a 3x3 grid.
Collect 3 dossiers in 30 turns.
Tools (use exactly one per turn):
TOOL: move(direction="north")
TOOL: talk(message="hello")
TOOL: collect()
Respond with exactly one TOOL: line.
Write one THOUGHT: line, then
one TOOL: line.

What the model wrote this step

TOOL: talk(message="Dr. Vapnik, I
seek the global minimum. Do you have
any jobs for an agent like me?")

gemini-2.5-flash, 16 September 2026

Yao and others, ReAct, ICLR 2023

Rajasekaran and others, Effective

context engineering, Anthropic, 2025

Where the observation comes from



```
TOOL: move(direction="east")
```

```
RESULT: You moved east to (2, 1).
```

```
Current location: A small beach. The sand is warm beneath your feet.
```

```
North: Open ground
```

```
South: Edge of the island (impassable).
```

```
East: Jungle
```

```
West: The south shore.
```

```
In this cell: Informant
```

*A recording, gemini-2.5-flash,
16 September 2026*

*Yao and others, ReAct, ICLR 2023,
from the released code*

*Hugging Face Agents Course,
Unit 1*

Where the observation comes from



```
TOOL: move(direction="east")
```

RESULT: You moved east to (2, 1).
Current location: A small beach. The sand is warm beneath your feet.
North: Open ground
South: Edge of the island (impassable).
East: Jungle
West: The south shore.
In this cell: Informant

Nobody
ran this

*A recording, gemini-2.5-flash,
16 September 2026*

*Yao and others, ReAct, ICLR 2023,
from the released code*

*Hugging Face Agents Course,
Unit 1*

Where the observation comes from



TOOL: move(direction="east")

The model kept writing.

RESULT: You moved east to (2, 1).
Current location: A small beach. The sand is warm beneath your feet.
North: Open ground
South: Edge of the island (impassable).
East: Jungle
West: The south shore.
In this cell: Informant

Nobody
ran this

*A recording, gemini-2.5-flash,
16 September 2026*

*Yao and others, ReAct, ICLR 2023,
from the released code*

*Hugging Face Agents Course,
Unit 1*

Where the observation comes from



```
TOOL: move(direction="east")
```

The model kept writing.


RESULT: You moved east to (2, 1).
Current location: A small beach. The sand is warm beneath your feet.
North: Open ground
South: Edge of the island (impassable).
East: Jungle
West: The south shore.
In this cell: Informant

Nobody
ran this

*A recording, gemini-2.5-flash,
16 September 2026*

*Yao and others, ReAct, ICLR 2023,
from the released code*

*Hugging Face Agents Course,
Unit 1*



Where the observation comes from



```
TOOL: move(direction="east")
```

The model kept writing.


RESULT: You moved east to (2, 1).
Current location: A small beach. The sand is warm beneath your feet.
North: Open ground
South: Edge of the island (impassable).
East: Jungle
West: The south shore.
In this cell: Informant

Nobody
ran this

*A recording, gemini-2.5-flash,
16 September 2026*

*Yao and others, ReAct, ICLR 2023,
from the released code*

*Hugging Face Agents Course,
Unit 1*



Where the observation comes from



```
TOOL: move(direction="east")
```

RESULT: You moved east to (2, 1).
Current location: A small beach. The sand is warm beneath your feet.
North: Open ground
South: Edge of the island (impassable).
East: Jungle
West: The south shore.
In this cell: Informant

Nobody
ran this


The model kept writing.

The harness kept the first call
and threw the rest away.

*A recording, gemini-2.5-flash,
16 September 2026*

*Yao and others, ReAct, ICLR 2023,
from the released code*

*Hugging Face Agents Course,
Unit 1*



Where the observation comes from




```
TOOL: move(direction="east")
```

RESULT: You moved east to (2, 1).
Current location: A small beach. The sand is warm beneath your feet.
North: Open ground
South: Edge of the island (impassable).
East: Jungle
West: The south shore.
In this cell: Informant

Nobody
ran this

Stop here

A regular expression, then move("east")



The model kept writing.

The harness kept the first call
and threw the rest away.

ReAct stopped generation at the
word Observation.

*A recording, gemini-2.5-flash,
16 September 2026*

*Yao and others, ReAct, ICLR 2023,
from the released code*


*Hugging Face Agents Course,
Unit 1*

3

Section 3


One call, three spellings

Who runs the tool




*Anthropic, Claude API
documentation, tool use,
2026*

Who runs the tool




*Anthropic, Claude API
documentation, tool use,
2026*

Who runs the tool




*Anthropic, Claude API
documentation, tool use,
2026*

Who runs the tool



*Anthropic, Claude API
documentation, tool use,
2026*

Who runs the tool




Who executed that call?

The harness

Where did the answer come from?

Anthropic, Claude API documentation, tool use, 2026

Who runs the tool



Who executed that call?

The harness

Where did the answer come from?

The world, through a tool

Anthropic, Claude API documentation, tool use, 2026

One tool call, written as ReAct text

The model wrote one line. The game answered.



What the model wrote

```
TOOL: move(direction="north")
```

*Yao and others, ReAct, ICLR
2023*

One tool call, written as ReAct text

The model wrote one line. The game answered.



What the model wrote

A format described in the
prompt

*Yao and others, ReAct, ICLR
2023*

```
TOOL: move(direction="north")
```

One tool call, written as ReAct text

The model wrote one line. The game answered.



What the model wrote

A format described in the
prompt

*Yao and others, ReAct, ICLR
2023*

```
TOOL: move(direction="north")
```

What the harness did

```
A regular expression, then move("north")
```

One tool call, written as ReAct text

The model wrote one line. The game answered.



What the model wrote

A format described in the
prompt

*Yao and others, ReAct, ICLR
2023*

```
TOOL: move(direction="north")
```

What the harness did

```
A regular expression, then move("north")
```

What came back

```
Moved north to (1, 0).
```

One tool call, written as ReAct text

The model wrote one line. The game answered.



What the model wrote

A format described in the
prompt

*Yao and others, ReAct, ICLR
2023*

```
TOOL: move(direction="north")
```

What the harness did

```
A regular expression, then move("north")
```

What came back

The same answer on all three
of the next frames

```
Moved north to (1, 0).
```

One tool call, written as ~~ReAct~~text a native tool call



What the model wrote

```
{ "name": "move",  
  "description": "Move one step north,  
                  south, east or west.",  
  "properties": { "direction": "string" },  
  "required": [ "direction" ] }
```

*Anthropic, Claude API
documentation, tool use, 2026*

One tool call, written as ~~ReAct~~text a native tool call



What the model wrote

- 1 The name the model will write

```
{ "name": "move",  
  "description": "Move one step north,  
                  south, east or west.",  
  "properties": { "direction": "string" },  
  "required": [ "direction" ] }
```

*Anthropic, Claude API
documentation, tool use, 2026*

One tool call, written as ~~ReAct~~text a native tool call



What the model wrote

- 1 The name the model will write
- 2 The sentence the model reads to decide

```
{ "name": "move",  
  "description": "Move one step north,  
                  south, east or west.",  
  "properties": { "direction": "string" },  
  "required": [ "direction" ] }
```

*Anthropic, Claude API
documentation, tool use, 2026*

One tool call, written as ~~ReAct~~text a native tool call



What the model wrote

- ① The name the model will write
- ② The sentence the model reads to decide
- ③ Sent on every request, and billed as input tokens

A few hundred tokens on every request, before your own prompt

```
{ "name": "move",  
  "description": "Move one step north,  
                  south, east or west.",  
  "properties": { "direction": "string" },  
  "required": [ "direction" ] }
```

*Anthropic, Claude API
documentation, tool use, 2026*

One tool call, written as ~~ReAct text~~ a native tool call



What the model wrote

```
{"name": "move", "description": ...,  
  "properties": ..., "required": ...}
```

```
{"type": "tool_use", "id": "toolu_01" ,  
  "name": "move",  
  "input": {"direction": "north"}}
```

*Anthropic, Claude API
documentation, tool use, 2026*

One tool call, written as ~~ReAct~~ text a native tool call



What the model wrote

```
{"name": "move", "description": ...,  
  "properties": ..., "required": ...}
```

```
{"type": "tool_use", "id": "toolu_01",  
  "name": "move",  
  "input": {"direction": "north"}}
```

What the harness did

The harness looks up "move" and runs it

*Anthropic, Claude API
documentation, tool use, 2026*

One tool call, written as ~~ReAct~~ text a native tool call



What the model wrote

```
{"name": "move", "description": ...,  
  "properties": ..., "required": ...}
```

```
{"type": "tool_use", "id": "toolu_01",  
  "name": "move",  
  "input": {"direction": "north"}}
```

What the harness did

The harness looks up "move" and runs it

What came back

*Anthropic, Claude API
documentation, tool use, 2026*

```
{"type": "tool_result",  
  "tool_use_id": "toolu_01", "content":  
  "Moved north to (1, 0)."} 
```

One tool call, written as ~~a native tool call~~ tokens

The same boxes as this morning. This one is not a word.



What the model wrote

<|channel|>

commentary


to=functions.move

One tool call, written as ~~a native tool call~~ tokens

The same boxes as this morning. This one is not a word.



What the model wrote




One tool call, written as ~~a native tool call~~ tokens

The same boxes as this morning. This one is not a word.



What the model wrote



One tool call, written as ~~a native tool call~~ tokens

The same boxes as this morning. This one is not a word.



What the model wrote

```
<|start|>assistant<|channel|>  
commentary  
to=functions.get_current_weather  
<|constrain|>json<|message|>  
{"location":"San Francisco"}<|call|>
```

<|channel|>

commentary

to=functions.move

Generation
stops here

<|message|>

{"direction":"north"}

<|call|>

*OpenAI, Harmony response
format, OpenAI Cookbook, 2025*

One tool call, written as ~~a native tool call~~ tokens

The same boxes as this morning. This one is not a word.



What the model wrote

```
<|start|>assistant<|channel|>  
commentary  
to=functions.get_current_weather  
<|constrain|>json<|message|>  
{"location":"San Francisco"}<|call|>
```

<|channel|>

commentary

to=functions.move

Generation
stops here

<|message|>

{"direction":"north"}

<|call|>

What the harness did

The harness reads the call and runs move("north")

*OpenAI, Harmony response
format, OpenAI Cookbook, 2025*

One tool call, written as ~~a native tool call~~ tokens

The same boxes as this morning. This one is not a word.



What the model wrote

```
<|start|>assistant<|channel|>  
commentary  
to=functions.get_current_weather  
<|constrain|>json<|message|>  
{"location":"San Francisco"}<|call|>
```

<|channel|> commentary to=functions.move

Generation
stops here

<|message|> {"direction":"north"} <|call|>

What the harness did

The harness reads the call and runs move("north")

What came back

*OpenAI, Harmony response
format, OpenAI Cookbook, 2025*

```
<|start|>functions.move to=assistant<|channel|>commentary  
<|message|>
```

One tool call, written as ~~a native tool call~~ tokens

The same boxes as this morning. This one is not a word.



What the model wrote

```
<|start|>assistant<|channel|>
commentary
to=functions.get_current_weather
<|constrain|>json<|message|>
{"location":"San Francisco"}<|call|>
```

<|channel|> commentary to=functions.move

Generation
stops here

<|message|> {"direction":"north"} <|call|>

What the harness did

The harness reads the call and runs move("north")

What came back

*OpenAI, Harmony response
format, OpenAI Cookbook, 2025*

```
<|start|>functions.move to=assistant<|channel|>commentary
<|message|>
Moved north to (1, 0). <|end|>
```


4

Section 4

The Model Context Protocol

What MCP is


An open protocol: a standard plug between an application that uses a model and the servers that offer it tools and data.



Model Context Protocol, Specification, 2026

What MCP is


An open protocol: a standard plug between an application that uses a model and the servers that offer it tools and data.



Model Context Protocol, Specification, 2026

What MCP is

An open protocol: a standard plug between an application that uses a model and the servers that offer it tools and data.



Tools: functions the model may ask for, run by the server


Resources: context and data, for the user or the model

Prompts: templated messages and workflows for the user

Model Context Protocol, Specification, 2026


How a tool call travels through MCP

Model Context Protocol, Specification, 2026




How a tool call travels through MCP

Model Context Protocol, Specification, 2026




How a tool call travels through MCP

Model Context Protocol, Specification, 2026




How a tool call travels through MCP

Model Context Protocol, Specification, 2026




How a tool call travels through MCP

Model Context Protocol, Specification, 2026




How a tool call travels through MCP

Model Context Protocol, Specification, 2026



How a tool call travels through MCP

Model Context Protocol, Specification, 2026



5


Section 5

The harness

An agent is a system




Cherny, Claude Code best practices, Anthropic, 2025
Anthropic, Claude API documentation, tool use, 2026



An agent is a system




Cherny, Claude Code best practices, Anthropic, 2025
Anthropic, Claude API documentation, tool use, 2026



An agent is a system




Cherny, Claude Code best practices, Anthropic, 2025
Anthropic, Claude API documentation, tool use, 2026



An agent is a system



Cherny, Claude Code best practices, Anthropic, 2025
Anthropic, Claude API documentation, tool use, 2026



Play the spy game



https://colab.research.google.com/github/eth-bmai-hs26/w2-cx-public/blob/main/spy-game/the_hidden_layer_micro_mission.ipynb

- 1 Open the notebook in Colab
- 2 Add your OpenRouter key as a Colab secret named OPENROUTER_API_KEY
- 3 Run the micro mission

References: definitions and the loop

Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*, 2nd edition. MIT Press, 2018.
incompleteideas.net/book/the-book-2nd.html

Erik Schluntz and Barry Zhang. Building effective agents. Anthropic, Engineering, 2024. anthropic.com/engineering/building-effective-agents

Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan and Yuan Cao. ReAct: synergizing reasoning and acting in language models. *ICLR 2023*, oral. arxiv.org/abs/2210.03629

Hugging Face. Agents Course, Unit 1, “Actions” and “Dummy agent library”. huggingface.co/learn/agents-course/unit1/actions

Prithvi Rajasekaran, Ethan Dixon, Carly Ryan and Jeremy Hadfield. Effective context engineering for AI agents. Anthropic, Engineering, 2025.
anthropic.com/engineering/effective-context-engineering-for-ai-agents

OpenAI. Running agents, Agents SDK documentation. openai.github.io/openai-agents-python/running_agents/

References: tool calling and the harness

Atty Eleti, Jeff Harris and Logan Kilpatrick. Function calling and other API updates. OpenAI, 13 June 2023. openai.com/index/function-calling-and-other-api-updates/

OpenAI. Function calling guide, API documentation. developers.openai.com/api/docs/guides/function-calling

Anthropic. Tool use: overview, define tools, handle tool calls, how tool use works. Claude API documentation. platform.claude.com/docs/en/agents-and-tools/tool-use/overview

OpenAI. Harmony response format. OpenAI Cookbook, 5 August 2025. developers.openai.com/cookbook/articles/openai-harmony

Model Context Protocol. Specification. modelcontextprotocol.io/specification/latest

Justin Young. Effective harnesses for long-running agents. Anthropic, Engineering, 26 November 2025. anthropic.com/engineering/effective-harnesses-for-long-running-agents

Boris Cherny. Claude Code: best practices for agentic coding. Anthropic, Engineering, 18 April 2025. web.archive.org/web/20250420034405/anthropic.com/engineering/claude-code-best-practices